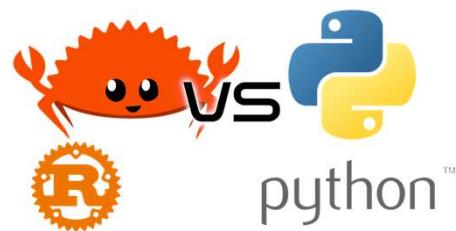




Quantitative Comparison

Who wins in the race of time and space?



17th October 2025

Contents

01 Introduction

02 Experiment Setup

03 Observations & Results - Preprocessing

04 Observations & Results - Searching

05 Conclusions & Next steps

Introduction

Quantitatively assess who does better in the race of time and space.

Python

- Interpreted Language
- Ease of Use
- **Memory Management** (Uses automatic garbage collection)
- Exceptionally rich library support for **data science**
- Faster development and experimentation cycles
- Parallelism(Constrained by GIL)
- Error Handling

Rust

- Compiler Language
- Suitable for performance-critical applications
- Uses an **ownership model** for safe and deterministic memory handling.
- Enforces strong compile-time checks
- Concurrency
- Growing ecosystem (notably **Polars**, **ndarray**, **Arrow** for data processing).
- Best suited for high-performance

Experimental Setup

We are assessing performance and memory in two scenarios

Preprocessing

- Detect Column types
- Missing value detection
- Imputation
- Normalization
- Add / Drop Columns
- Filter rows
- Sorting
- Sampling

Searching Algorithms

- Linear Search
- Binary Search
- Interpolation Search
- Jump Search

Preprocessing

Dataset size

- Healthcare claims dataset
- Size of Data : (116352, 32)
- ~14MB in size

Searching Algorithms

Array of size 1 million within a range MIN and MAX

Statistics obtained on scenarios searching

- First
- Middle
- Last
- Element < MIN
- Element > MAX

Time profiling

- Python - line profiler
- Rust - Measuring Elapsed time
- Hyperfine - Benchmarking for execution time

Memory Profiling

- Python – memory_profiler
- Rust – using sysinfo



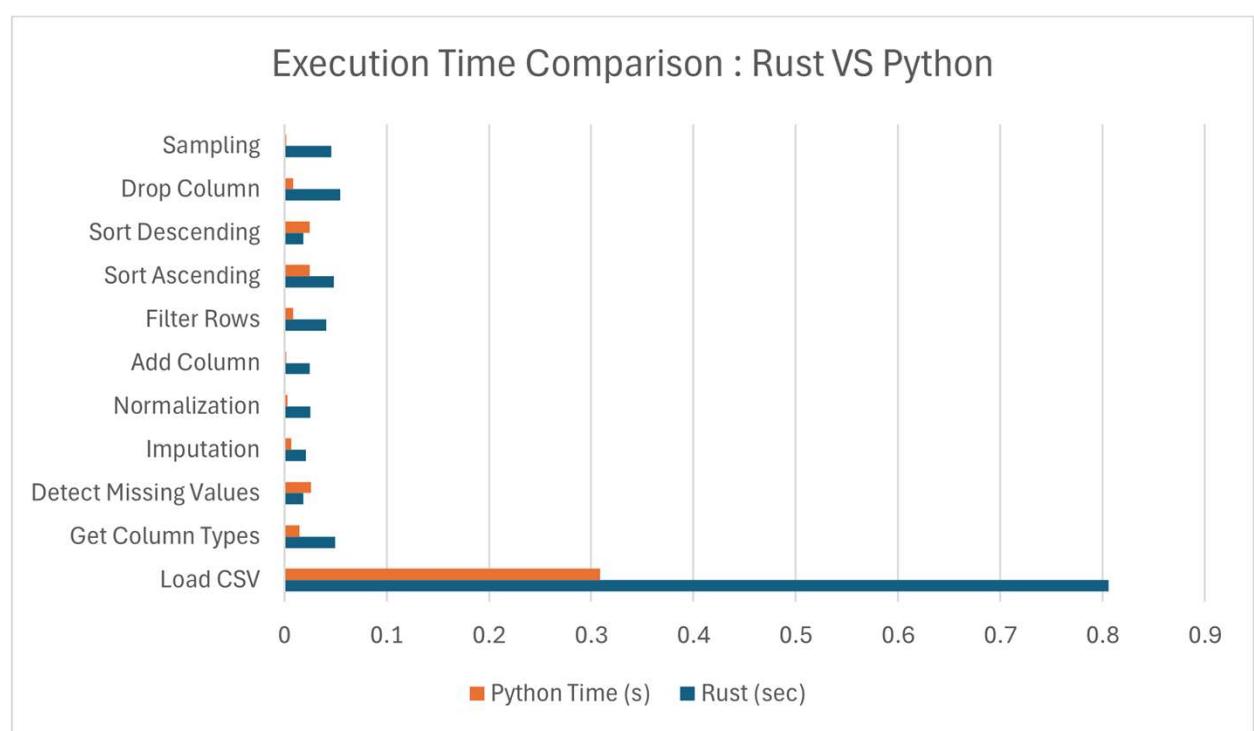
Observations & Results

Preprocessing

Time Profiling Comparison

Python wins at each individual task execution time

- Loading dominated in both
- Python dominates in execution time per operation
- Visualizations shows Python's optimization and efficiency in data operations



Benchmarking tells us a different story

Warmup Impact

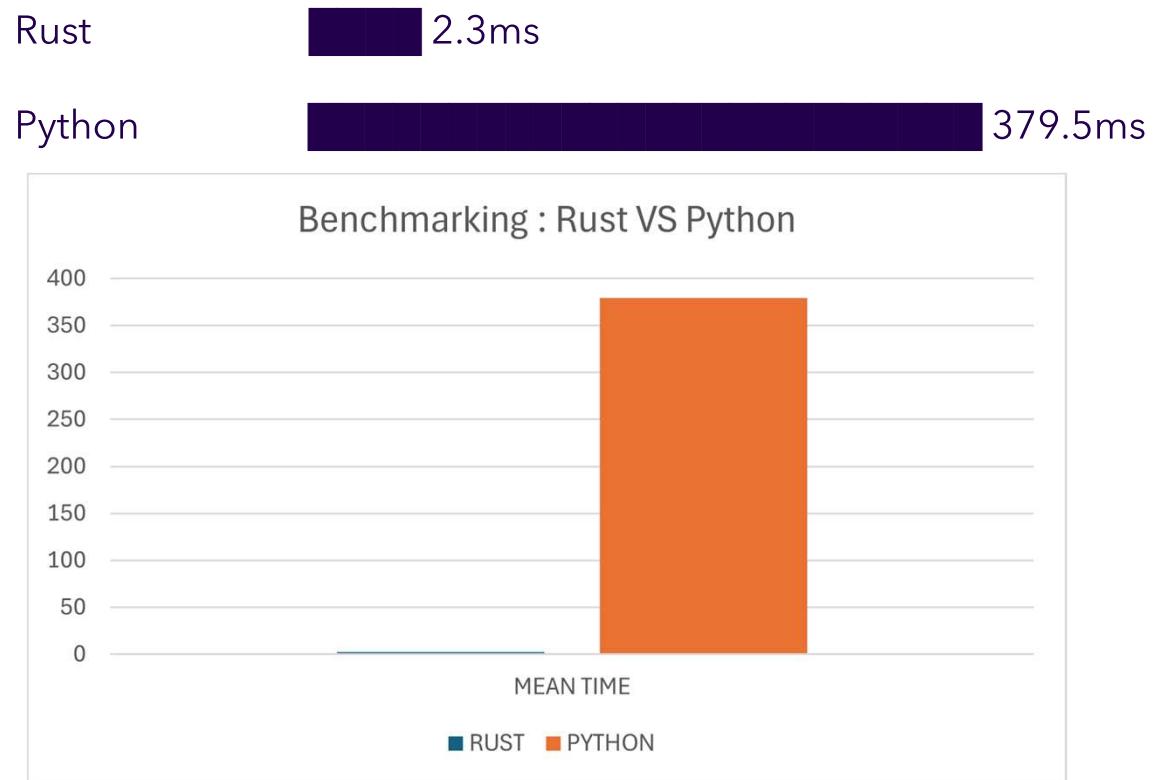
Stability

- Warmup (3 runs): Rust stabilized instantly, Python took ~0.39s.
- Warmup (10 runs): Rust stable at 0.00s, Python stable around 0.38s.

Improvement due to warmup

- Rust : 44%
- Python : 4%

This shows that the reason for slow per-operation execution is due to cache misses in course of process execution



Memory Profiling

Rust consistently uses less memory than Python

Average memory footprint

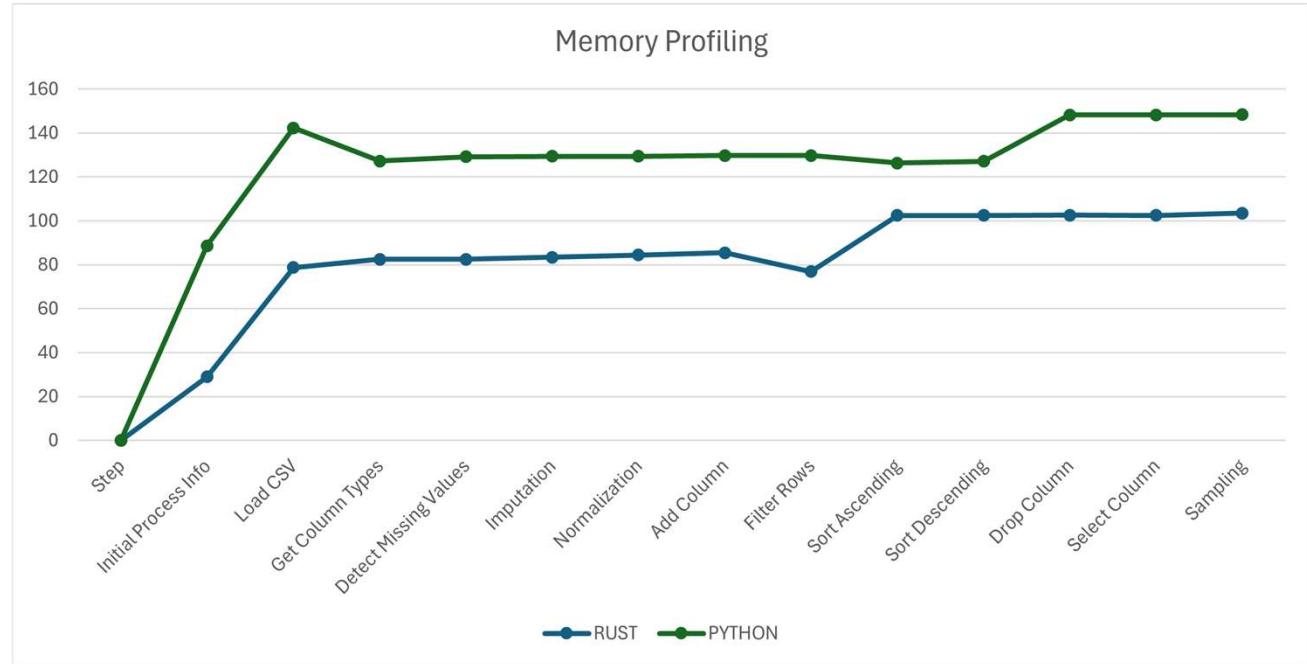
- Rust ~88 MB
- Python ~134 MB

Peak Memory

- Rust : 103MB
- Python : 148MB

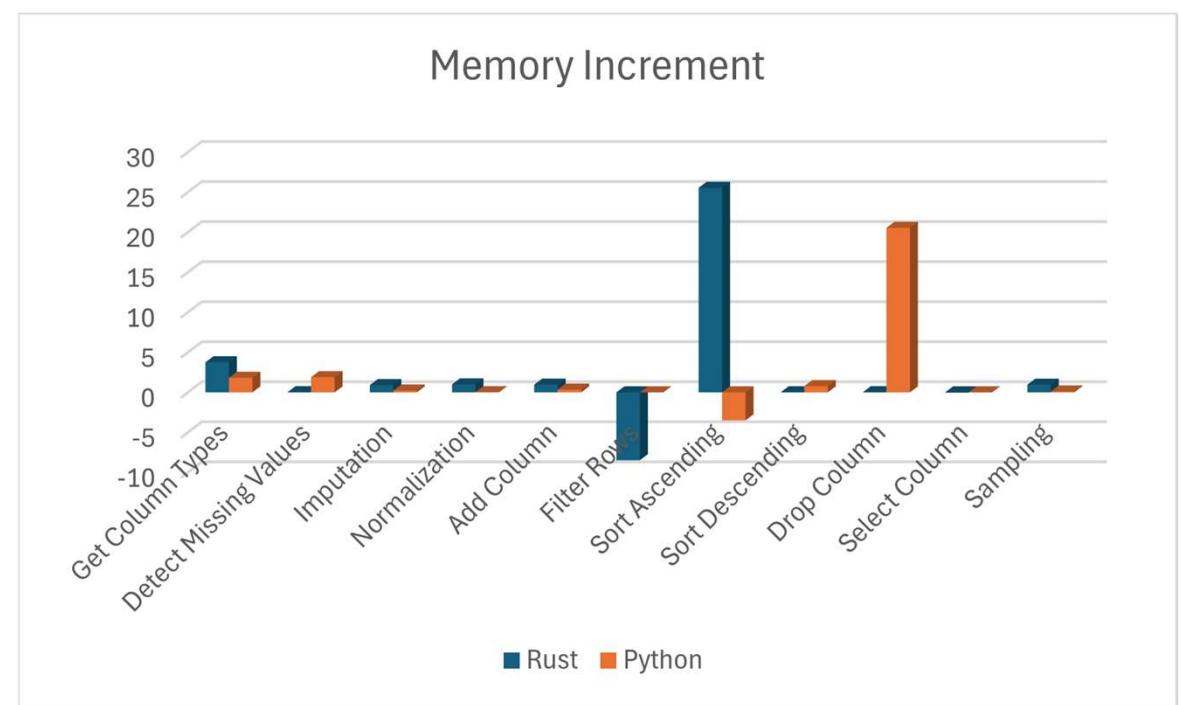
What is the use ?

More predictable **scaling for large datasets.**



Rust shows controlled memory reuse (drops after filtering), while Python retains buffers.

- Could be the result of selected rows getting dereferenced.(Ownership & Borrowing)





Observations & Results

Searching Algorithms

Time Profiling

We observe that Python performs significantly better than Rust

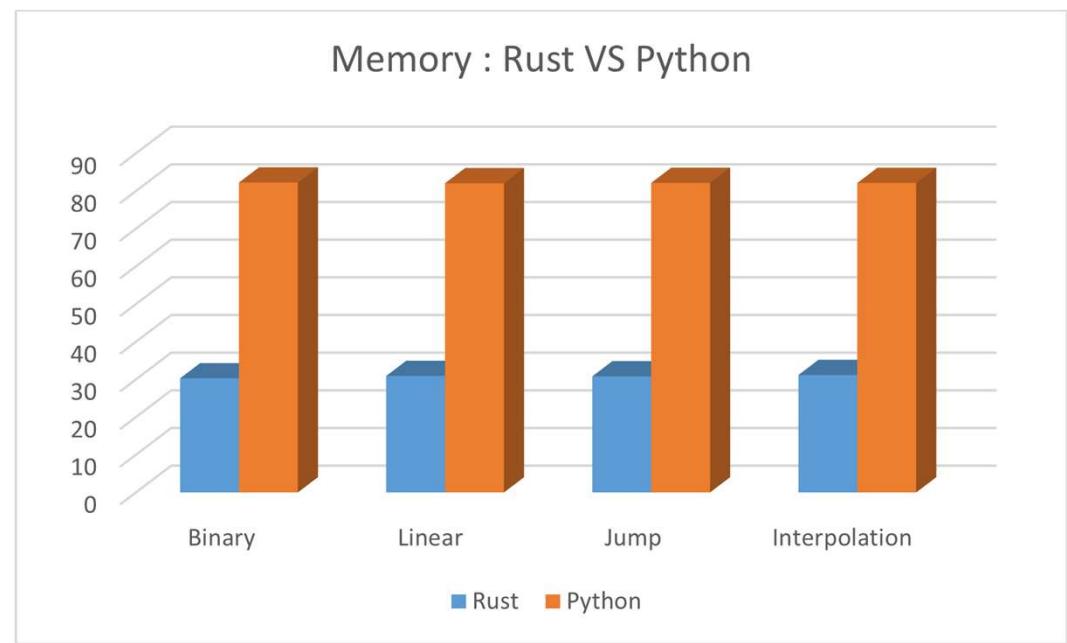
These are results of benchmarking on hyperfine

- There is no help even after warmup

	Rust	Python
Search Algorithm	Time sec	Time sec
Binary	1.124155	0.26
Interpolation	1.07282	0.32126
Jump	1.024465	0.321994
Linear	1.041261	1.04

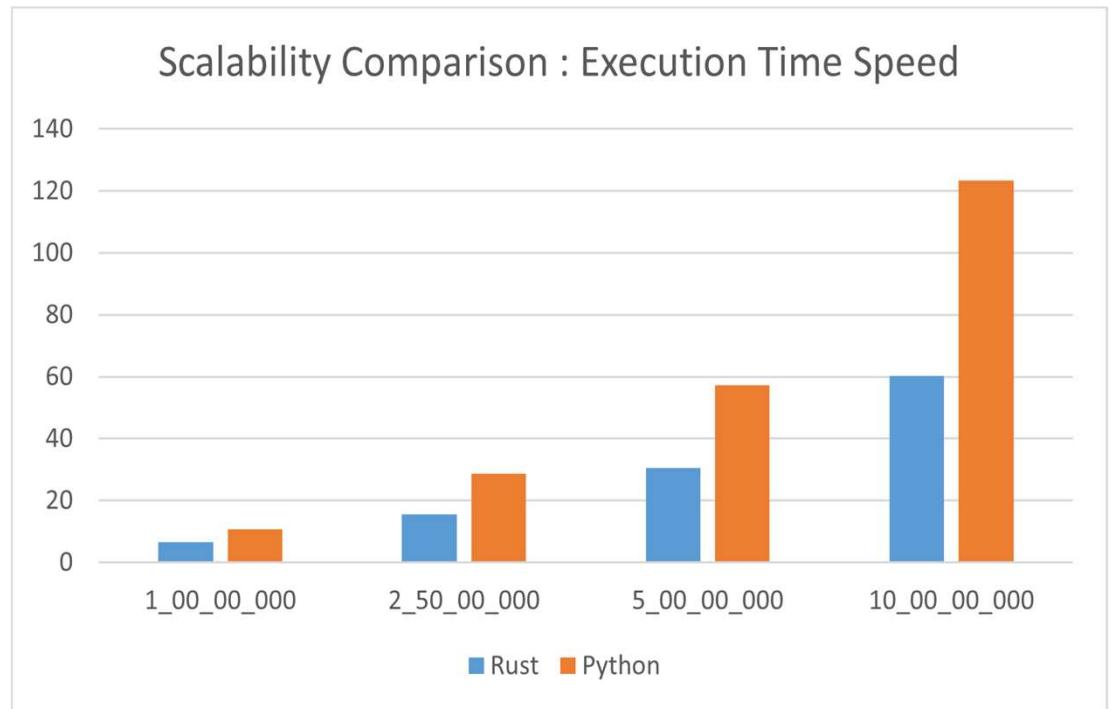
Memory Profiling

- Rust is much more memory-efficient
 - roughly **2.7× less memory** than Python across all algorithms
- The **difference across algorithms** within the same language is minimal
- The **variation** (max-min) within:
 - Rust: ~0.9 MB (very stable)
 - Python: ~0.2 MB (also stable)



Scalability

- What happens when we increase array size?
- Who scales better?
- Consider the search algorithm with max time complexity
- The comparative study shows that **Rust consistently outperforms Python** in both **execution time** and **memory efficiency**



Conclusions

- Time Performance
 - A. Rust outperforms Python in both preprocessing and searching speed.
 - B. Shows better scalability and is ideal for high-performance ETL and real-time data processing.
 - C. Python remains stronger for rapid prototyping and flexibility.
- Memory Performance
 - A. Rust uses less and more predictable memory across all tasks.
 - B. Its ownership model ensures efficient reuse, unlike Python's garbage collection.
 - C. Rust offers stable, scalable performance, while Python suits analytical experimentation.

Next Steps

- Integrate **multi-threading and parallelism** in Rust and **multiprocessing in Python** to evaluate performance gains and CPU utilization under concurrent workloads.
- Expand the scope to other algorithmic domains such as **sorting, regression, and clustering** to generalize the performance trends observed.





Pranav Phanindra Sai Manepalli

pranavphanindrasai.manepalli@zelis.com

7981145597

The information in this document is for general informational purposes only on an "as is" basis. Zelis makes no representations or warranties of any kind, regarding this information, its products and/or services featured herein, or any specific outcomes based upon the use of such information/products. Nothing contained in this document is intended to constitute legal advice. With respect to services provided by our Price Optimization business unit, Zelis provides savings and coding recommendations for each client, who independently assess whether to accept them. All information related to a client's approach to its cost containment program is confidential.

